

Evaluation report - SerafinoSignal

Answered Problem Sections:

4.5 / 5

My core and bonus request were both answered in this solution, which were recurring complaint themes across restaurants and patterns by restaurant types. The operational themes were clearly separated from the rest. The only thing/ what the .5 point is because some of the 11 clusters might not be considered complaints (ex. positive sentiments, food type mentions like 'pizza experience').

Technique Judgement:

5 / 5

I would say picking unsupervised clustering here was the right call. I never gave them labeled complaint categories so there wasn't anything to supervise. They also didn't just pick one method, they tried multiple (TF-IDF, sentence embeddings, and UMAP+HDBSCAN). They also explained why the embedding-based KMeans won which shows that they understood the problem and tested many options for the best solution.

Evaluation/Interpretation Quality:

5 / 5

The interpretation is tied back to what I actually care about, which is the restaurant coverage, avg rating per theme, and a 'lift vs platform average' comparison that shows which restaurant types over-index on which complaints.

Communication:

5 / 5

The parts that are meant for the stakeholder (findings, summary, conclusion) are all written in easy to understand language. The streamlit app also offers a very adjustable, exploratory way of looking at the findings and data.

Open Feedback:

This analysis was very thorough and honest. The model choices are defended and the limitations are stated clearly. The four operational themes it landed on are definitely the type of thing an outreach director could use/act on. The findings are what I asked for, it includes recurring and actionable operational themes with patterns across different restaurants. The dashboard is very clean and easy to navigate.